

Python Bootcamp 2025-26

(Statistics and Data Analysis course)

Teachers: [Jorge Carretero](#), [Pau Tallada](#) & [Francesc Torradeflot](#)

<https://mastercosmosbcn.cat/>

Prerequisites

- Laptop with:
 - Anaconda (~5GB of storage!)
 - Python environment with the following libraries installed:
 - NumPy • SciPy • Pandas • Astropy • Matplotlib • Scikit-learn • (PyTorch)
 - Git
 - VSCode
- For communicating during classes:
 - Connection details: <https://meet.google.com/qgt-gdwb-nqt>

General information

- [Master official information](#)
- In the Virtual Campus: *Estadística i Anàlisi de Dades* [MO79240]
 - **Part 0: Python transversal course**
- [Master Schedule](#)
- 6 classes; 2 hours each, always from 10:00 - 12:00
 - **Sep 23, 26, 30, Oct 3, 7, 10**
- Classroom: C7/029
- Contact:
 - Jorge Carretero (carretero@pic.es) - Physicists
 - Pau Tallada (tallada@pic.es) - Research Software Engineer (RSE)
 - Francesc Torradeflot (torradeflot@pic.es) - Mathematician & RSE

Course structure

- **Project-oriented course (group work)**
 - Includes some theoretical sections
- Main goal: **Train a machine learning algorithm using images**
- Sessions:
 - Session 1 (23 Sept): Introduction and Setup
 - Session 2 (26 Sept): Data Exploration
 - Session 3 (30 Sept): Plotting
 - Session 4 (3 Oct): Advanced Programming Concepts
 - Session 5 (7 Oct): Miscellaneous Topics
 - Session 6 (10 Oct): Project Presentation & Feedback

Evaluation Criteria

- Active participation during the course (80%):
 - Attending the sessions
 - Presenting group work
 - At least 3 groups each week
 - Asking and answering questions
- Effort (20%):
 - Going beyond the suggested implementations
 - Proposing innovative solutions
 - Other significant contributions
- Overall grade: 10% of the total mark of the course

Practical note: Post-its will be given to people who contribute; they should be filled in with your name(s)

Today's course

- Quick self-introductions (very very brief, please!)
 - Example: "I'm Jorge Carretero, a physicist and I'm interested on galaxy surveys."
- Project Overview and Introduction
- Last part:
 - Confirm that everyone has the proper setup
 - **Define the configuration of the different groups (3 people per group)**
 - Can also be done during the mid-course break
 - Propose the first exercise for the next session
- Open Meet connection to be able to share screens:
 - Video call link: <https://meet.google.com/qgt-gdwb-nqt>

Welcome & Overview (I)

- Purpose of the Bootcamp
 - Equip you with **modern programming tools used in research and industry**
 - Build habits for **reproducibility, collaboration, and scalability**
- Why Python?
 - Ubiquitous in science, data analysis, and AI
 - Huge ecosystem of open-source libraries for astronomy, simulation, and ML
 - Easy to learn, readable, and powerful enough for large-scale computation

Welcome & Overview (II)

- Learning Outcomes

- Set up a productive development environment (IDEs, notebooks)
- Use Git & GitHub for version control and collaboration
- Automate workflows with CI/CD
- Write reliable code through testing & TDD
- Leverage LLM coding assistants (Copilot, Gemini)

- Course Themes

- Reproducibility: code & data pipelines that others can trust
- Collaboration: working together effectively with modern tools
- Productivity: using automation and best practices to move faster

Python in Scientific Computing

- Why Python Is the Researcher's Language
 - Ecosystem for Science
 - NumPy • SciPy • Pandas • Astropy • Matplotlib • PyTorch • Scikit-learn
 - Performance
 - Interoperates with C/Fortran via Numba, Cython
 - Community
 - Open-source, peer-reviewed, collaborative culture
 - Adoption
 - Standard in astronomy, cosmology, and data science

Intro to Programming Languages

What Is a Programming Language?

- **Purpose:** Tool to give precise instructions to a computer
- **Levels:**
 - High-level: Python, R, Julia – human-readable
 - Low-level: C, Fortran, Assembly – close to hardware
- **Domain-Specific:** IDL, Mathematica – specialized for a field

Compiled vs Interpreted Languages

How Code Becomes Execution

- **Compiled:** C, Fortran — translated to machine code before running → very fast
- **Interpreted:** Python, R — executed line by line → easier to prototype
- **Performance Tradeoff:** Fast runtime vs. fast development cycle
- **JIT Compilation:** Numba, PyPy bring compiled speed to Python

IDEs & Development Environments

- **IDE** = Integrated Development Environment
 - VS Code, PyCharm — autocomplete, linting, debugging, testing, tooling
- **JupyterLab**: Interactive notebooks for science & exploration
- **Terminal** Tools: Run quick scripts, manage environments
- Recommended Setup:
 - Python 3.12+, VS Code, Jupyter, Git, venv/Conda

Jupyter Notebooks for Research

- Literate Programming
 - Combine code, text (Markdown), LaTeX, plots
- Use Cases
 - Data exploration, simulations, reproducible workflows
- Sharing
 - Easy to save, share, and rerun experiments
- Tools
 - JupyterLab • VS Code notebooks • Google Colab

Version Control with Git

- Version Control
 - Save history, revert changes, experiment safely
- Core Commands
 - init, clone, commit, push, pull, branch
- Collaboration
 - GitHub / GitLab for hosting & teamwork
- Best Practices
 - Small commits
 - Clear messages
 - .gitignore

Teamwork and Open Science

- Collaboration Tools:
 - Pull Requests • Code Reviews • GitHub Issues/Discussions
- Project Structure:
 - Modular code • Clear directories • Inline documentation
- Open Science & Licensing:
 - MIT, GPL, and other open-source licenses
- Best Practices:
 - Review carefully, communicate clearly, document changes

Agile & Scrum (for Scientists!)

Managing Research Like a Software Project

- Agile Mindset
 - Iterative, flexible, feedback-driven progress
- Scrum Basics
 - Sprints • Stand-ups • Retrospectives
- Adapted for Science
 - Weekly goals • Research backlogs • Milestone tracking
- Tools
 - Trello • GitHub Projects • Notion

Test-Driven Development (TDD)

Trust Your Code, Trust Your Results

- Why Test?
 - Catch bugs early, ensure reproducibility
- Test Types:
 - Unit • Integration • Regression
- TDD Workflow:
 - Write test → Write code → Validate
- Tools:
 - pytest • unittest • doctest

- **CI/CD Concepts:**
 - CI (Continuous Integration): Run tests on every code change
 - CD (Continuous Delivery/Deployment): Automate releases or reproducible pipelines
- **Automation Tasks**
 - Testing • Linting • Formatting
- **Tools**
 - GitHub Actions • GitLab CI • CircleCI
- **Example**
 - Python repo automatically runs tests and linters on each push

Package & Environment Management

- Why It Matters:
 - Reproducibility • Dependency isolation
- Tools & Approaches:
 - venv + pip — lightweight environment + package management
 - conda — science-focused packages, cross-platform support
- Tracking Dependencies:
 - requirements.txt • environment.yml • pyproject.toml
- Advanced Reproducibility:
 - Binder • Docker

Using LLMs & Co-Agents

- What They Do:
 - Suggest code, boilerplate, docstrings, test cases
- Use Cases:
 - Debugging, refactoring, algorithm ideas, documentation
- Tools:
 - GitHub Copilot • Gemini • ChatGPT
- Pitfalls:
 - Always review suggestions, **don't trust blindly**
- Ethics:
 - Maintain code quality, credit sources, ensure reproducibility

Scale Your Computations

- Why Parallelize?
 - Handle large datasets, run heavy simulations faster
- Tools:
 - Dask: Scales NumPy/pandas for parallel processing
 - Ray: Python-native distributed tasks
 - Spark / PySpark: Big data analytics across clusters
- Cluster Environments
 - HPC basics, SLURM, MPI, job scripts

Make Your Code Understandable and Reusable

- Why Document?
 - Ensure reproducibility • Shareable knowledge • Maintainable code
- Docstring Standards:
 - NumPy/SciPy style — describe inputs, outputs, purpose, examples
- Tools:
 - `help()` in Python • Sphinx • MkDocs for auto-generated documentation

Keep Your Code Clean and Consistent

- Linting:
 - flake8, pylint — detect errors, enforce style
- Formatting:
 - ruff, black, isort — automatic code formatting and import sorting
- Pre-commit Hooks:
 - Run checks before every commit
- Principle:
 - Write clean, readable code for yourself and collaborators

Summary & Setup Instructions

Recap & Get Started

- Recap of Key Tools:
 - IDEs, Git, Testing, CI/CD, Notebooks, LLMs
- Recommended Setup:
 - Python 3.12+, VS Code, JupyterLab, Git, pytest
- Support & Resources
 - GitHub repo, Slack/Discord channels, install guides
- Next Steps
 - Build your development environment, follow setup instructions in shared repo

Configuration Setup – Troubleshooting

- Define the configuration of the different groups
- Verify that everyone has the correct setup (troubleshooting)
 - Identify common configuration issues

Exercise 1: Documentation & Data retrieval

- Data Sources – Galaxy Zoo 1
 - Where to access it: [SDSS Skyserver](#)
- For the next session, we need two datasets:
 - Galaxy morphology classifications (table "zooSpec")
 - Labels: elliptical or spiral (used as the target for training the ML model)
 - Photometry data for galaxies with morphology classifications (table "PhotoObjDR7")
 - Used to select objects of interest and as features for the ML model
- Hint – SQL query (JOIN between those two tables)

Exercise 1: Documentation & Data retrieval

- Data Sources – Galaxy Zoo 1
 - Where to access it: [SDSS Skyserver](#) (account needed!)
- For the next session, we need two datasets:
 - Galaxy morphology classifications (table "zooSpec")
 - Labels: elliptical or spiral (used as the target for training the ML model)
 - Photometry data for galaxies with morphology classifications (table "PhotoObjDR7")
 - Used to select objects of interest and as features for the ML model
- Hint – SQL query (JOIN between those two tables)
 - Submit a CAS job with the SQL query below and download in csv format:

```
SELECT ZooSpec.*, PhotoObjDR7.* into MyDB.ZooSpecPhoto
FROM ZooSpec INNER JOIN PhotoObjDR7
ON PhotoObjDR7.dr7objid = ZooSpec.dr7objid
```